



XV6 SCHEDULING

Operating Systems

University of Tehran-Faculty of Computer Engineering

Fall 98

LET'S SEE HOW XV6 DOES SCHEDULING

- Main.c → scheduler() / Proc.c → scheduler() → Round Robin Implementation

```
static void
mpmain(void)
{
    if(cpu() != mpbcpu())
        lapicinit(cpu());
    ksegment();
    cprintf("cpu%d: mpmain\n", cpu());
    idtinit();
    xchg(&c->booted, 1);

    cprintf("cpu%d: scheduling\n", cpu());
    scheduler();
}
```

```
void
scheduler(void)
{
    struct proc *p;

    for(;;){
        // Enable interrupts on this processor, in lieu of saving intena.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before running back to us.
            cp = p;
            usegment();
            p->state = RUNNING;
            swtch(&c->context, &p->context);

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            cp = 0;
            usegment();
        }
        release(&ptable.lock);
    }
}
```

LET'S SEE HOW XV6 DOES SCHEDULING: SCHEDULER() FUNCTION

- `swtch(&c->context, &p->context);`
 - Makes process "p" run in next time quantum by substituting context pointers
 - `&c->context`: pointer to current CPU scheduler context
 - `&p->context`: pointer to next running process context
- What happens if a process is paused by timer interrupt or is blocked by I/O operation?
- How can we pick another process to run?

LET'S SEE HOW XV6 DOES SCHEDULING: CHOOSING ANOTHER PROCESS

- Assume we have timer interrupt
 - Timer generates interrupt → Cause syscall to call a trap(implemented in trap.c)
 - Yield function is executed(implemented in proc.c)

```
// Force process to give up CPU on clock tick.  
// If interrupts were on while locks held, would need to check nlock.  
if(cp && cp->state == RUNNING && tf->trapno == T_IRQ0+IRQ_TIMER)  
    yield();
```

Trap.c

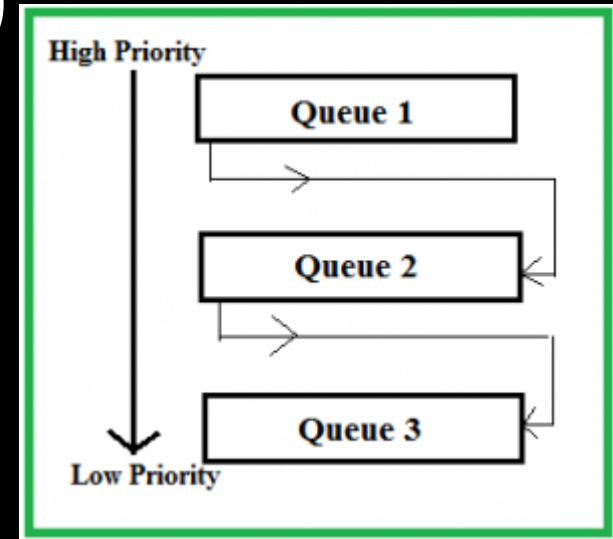
This mentioned procedure is implementation of
Round-Robin Scheduling in XV6!

```
// Give up the CPU for one scheduling round.  
void  
yield(void)  
{  
    acquire(&ptable.lock);  
    cp->state = RUNNABLE;  
    sched();  
    release(&ptable.lock);  
}
```

Proc.c

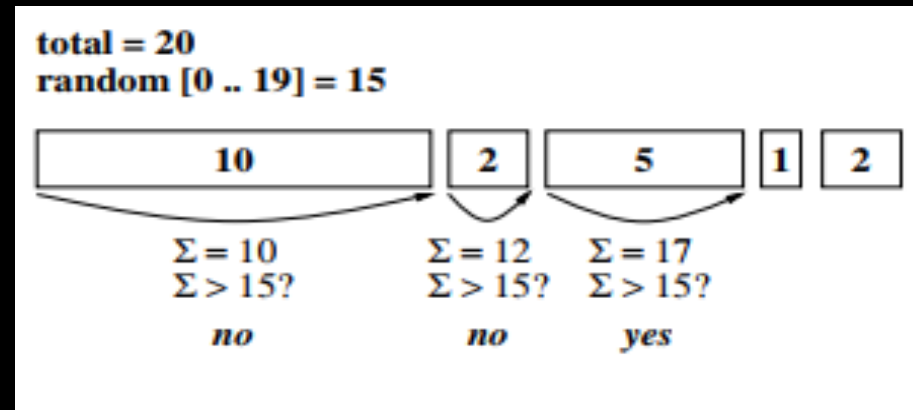
SCHEDULING MODIFICATIONS

- Round Robin → Multi Layer Scheduling
 - 1st level: Lottery Scheduling (First priority)
 - 2nd level: Highest Response Ratio Next (Second priority)
 - 3rd level: Shortest Remaining Priority First (Third priority)



LOTTERY SCHEDULING

- Generating lottery tickets for each process
- Generating random number in each interval and choose a process to run according to its lottery ticket numbers and generated random number
- Process with more lottery tickets is more probable to be chosen than process with less lottery tickets!



HRRN SCHEDULING

- You need to calculate a process waiting time and its executed cycles
- When a process executes, its executed cycle attributes increases 0.1 in magnitude, and the default value is set to 1
- The higher a process response ratio is, the higher the process chance is to be executed

$$HRRN = \frac{WaitingTime}{ExecutedCycleNumber}$$

$$WaitingTime = CurrentTime - ArrivalTime$$

SRPF SCHEDULING

- Every Process has an attribute named “Remaining Priority”
- Process with the lowest remaining priority is executed first
- When a process is executed, its remaining priority attribute will decrease 0.1 in magnitude
- **Attention!** The minimum magnitude of remaining priority is 0

COMPLEMENTARY SYSTEM CALLS

- 1. Change level of scheduling
- 2. Change remaining priority of processes in the last level
- 3. Assigning Lottery Tickets to 1st level processes
- 4. Listing all processes (helpful for your debugging)

Thank You for your
attendance 😊